


Week 1 – Stanford Machine Learning

Linear Regression

Let us create some random data, say housing sizes in square feet on x and price levels on y , and plot it.

In[823]:=

```
trainX = {100, 320, 213, 512, 58, 84, 113, 142, 93, 121, 421, 432, 249, 254};  
trainY = {140 000, 400 000, 241 000, 489 000, 78 000, 123 000,  
          139 000, 143 000, 97 000, 134 000, 392 000, 458 000, 311 000, 378 000};  
trainComposed = Transpose@{trainY, Table[1, {Length[trainX]]}, trainX};  
listPlot = ListPlot[Transpose@{trainX, trainY},  
                    AxesLabel → {"Size (feet2)", "Price ($)"}];  
Show[listPlot]
```



Now, we define a linear hypothesis function since our data seems to allow for a linear approximation

In[50]:=

```
Hyp[x_, θ_] := θ . x;
```

followed by the squared error cost function

In[51]:=

```
Cost[θ_] :=  $\frac{1}{2}$  Total[(Hyp[{##2}, θ] - #1)2 & @@@ trainComposed];
```

We then minimise the cost function, after which we plot the hypothesis function with the minimised values


In[885]:=

```

Clear[v]
θ = Array[v, Length[trainComposed[[1]]] - 1];
regression = Minimize[Cost[θ], θ]
regressionPlot := Plot[Hyp[{1, x}, {v[1], v[2]} /. regression[[2]]],
  {x, 0, Max[trainX]}, PlotStyle -> {Orange, Thin}]
Show[listPlot, regressionPlot]

```

Out[887]= $\left\{ \frac{10\,105\,881\,123\,250\,000}{998\,147}, \left\{ v[1] \rightarrow \frac{38\,758\,111\,500}{998\,147}, v[2] \rightarrow \frac{955\,610\,000}{998\,147} \right\} \right\}$



Gradient Descent

Let us repeat the regression analysis by using gradient descent instead, using the same data. First, we create the step function that returns the next step of the gradient descent algorithm using the update rule. Note that because of the huge values on the pricing axis, the learning rate parameters are particularly sensitive. Any learning rate above 1×10^{-7} will thus make the gradient descent algorithm overshoot. Instead of setting a very low learning rate parameter, as below, a better fix would be to represent the housing prices in terms of thousand dollars.

In[650]:=

```

Clear[v]
α = {0.1, 0.0000003};
θ = Array[v, Length[trainComposed[[1]]] - 1];
Step[θ_] := θ + α Sum[(trainY[[i]] - Hyp[trainComposed[[i, 2 ;;]], θ])
  trainComposed[[i, 2 ;;]], {i, 1, Length[trainX]}]

```


To see how we actually learn, we can make a contour plot of the cost function and of the results as we step through.

In[861]:=

```

contourPlot := ContourPlot[Cost[{x, y}],
  {x, 0, 400 000}, {y, 0, 980}, PlotLegends → Automatic];
gradientDescentPlot := ListLinePlot[NestList[Step, {1, 1}, 50],
  AxesLabel → {"x", "y"}, PlotStyle → {}]
Show[contourPlot, gradientDescentPlot]

```



Moreover, we can plot the step function as a function of the iterations to picture the learning rate. We do so by creating the learning function

In[690]:=

```

Learning[x_] := Cost@Nest[Step, {1, 1}, x];

```


We do indeed learn, as seen in the plot below

In[691]:=

```

learningRate = ListLinePlot[#, AxesLabel → {"Iterations", "Cost"}] &@
  Transpose@{Range[0, 50], Learning /@ Range[0, 50]}


```



We then step through the algorithm fifty times and plot all the lines of all intermediate steps to see how the regression converges towards the optimum

In[876]:=

```
gradientDescentRegressionPlot =
Show[Plot[Hyp[{1, x}, #], {x, 0, Max[trainX]}, PlotStyle -> {Orange, Thin}] & /@
(Nest[Step, {1, 1}, #] & /@ Range[0, 50])];
Show[listPlot, gradientDescentRegressionPlot]
```



Linear Regression using the Normal Equations

For data that only contains a small set of features we can easilier find the optimum by using the normal equations. Note that since these equations need to find the inverse, if the feature matrix is particularly large, it is better to use the gradient descent algorithm since finding the inverse is an $O(n^3)$ operation.

We prepare our data set

In[880]:=

```
X = Transpose@{Table[1, {Length[trainX]}], trainX};
Y = trainY;
```

and find the optimal parameters

In[882]:=

```
θ = Inverse[Transpose[X].X].Transpose[X].trainY;
```

Lastly, we plot the parameters to show that it is indeed the optimum.

In[883]:=

```
normalEquationPlot =
Plot[Hyp[{1, x}, θ], {x, 0, Max[trainX]}, PlotStyle -> {Orange, Thin}];
Show[listPlot, normalEquationPlot]
```

